

Python Numba for scientific code

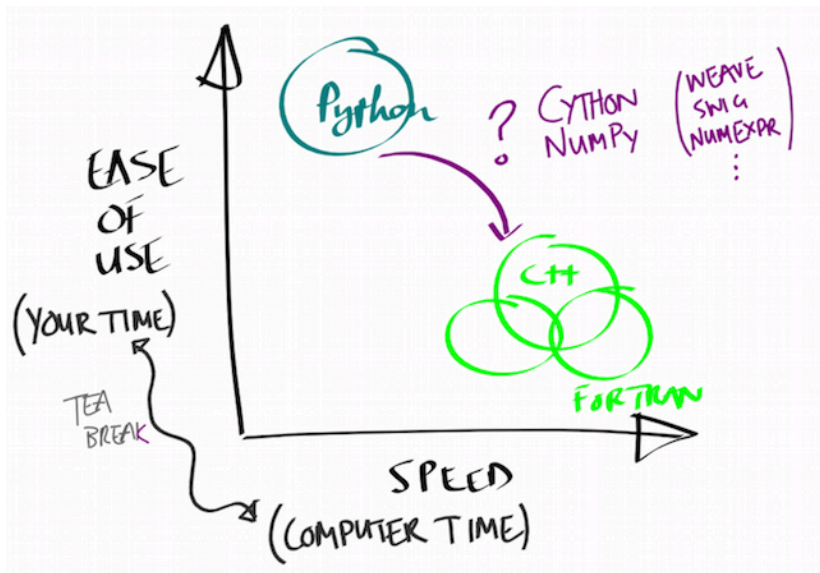
Zbigniew Jędrzejewski-Szmek

zbyszek@in.waw.pl



Instytut im. M. Nęckiego, 25.11.2018

Programmer time vs. computer time



Previous approaches

a.k.a. the graveyard of technologies

1. Use compiled C or C++ or Fortran code with cpython
 - ▶ C extension
 - ▶ .so library + ctypes
 - ▶ Fortran + f2py
 - ▶ boost-python
 - ▶ SWIG
 - ▶ (Cython)

Previous approaches

a.k.a. the graveyard of technologies

1. Use compiled C or C++ or Fortran code with cpython
 - ▶ C extension
 - ▶ .so library + ctypes
 - ▶ Fortran + f2py
 - ▶ boost-python
 - ▶ SWIG
 - ▶ (Cython)
2. A different Python interpreter
 - ▶ jython
 - ▶ psyco
 - ▶ pypy

Previous approaches

a.k.a. the graveyard of technologies

1. Use compiled C or C++ or Fortran code with cpython
 - ▶ C extension
 - ▶ .so library + ctypes
 - ▶ Fortran + f2py
 - ▶ boost-python
 - ▶ SWIG
 - ▶ (Cython)
2. A different Python interpreter
 - ▶ jython
 - ▶ psyco
 - ▶ pypy
3. Numpy

Previous approaches

a.k.a. the graveyard of technologies

1. Use compiled C or C++ or Fortran code with cpython
 - ▶ C extension
 - ▶ .so library + ctypes
 - ▶ Fortran + f2py
 - ▶ boost-python
 - ▶ SWIG
 - ▶ (Cython)
2. A different Python interpreter
 - ▶ jython
 - ▶ psyco
 - ▶ pypy
3. Numpy
4. “jit” just the inner loop
 - ▶ weave (Python + annotations to C++, inline)
 - ▶ pyrex (Python + annotations to C, external)
 - ▶ cython
 - ▶ numexpr (“jitting” of basic numpy array operations)
 - ▶ numba

Special syntax

numexpr

```
# numexpr_evaluate.py  
import numpy as np, numexpr as ne  
a = np.arange(1e6)  
b = np.random.randint(10, size=(1_000_000,))  
print(ne.evaluate('a*b-4.1*a > 2.5*b'))
```

Special syntax and type annotations

cython

```
# cython_integrate.pyx
def f(double x):
    y = (x*x*x - 3)*x
    return y
def integrate_f(double a, double b, int n):
    cdef:
        double dx = (b - a) / n
        double dx2 = dx / 2
        double s = f(a) * dx2
        int i = 0
    for i in range(1, n):
        s += f(a + i * dx) * dx
    s += f(b) * dx2
    return s
```


“jit”?

```
>>> import numba
```

```
>>> @numba.jit
```

```
... def f(x):
```

```
...     y = x*5 + x
```

```
...     return y
```

“jit”?

```
>>> import numba
```

```
>>> @numba.jit
```

```
... def f(x):
```

```
...     y = x*5 + x
```

```
...     return y
```

```
>>> f(1)
```

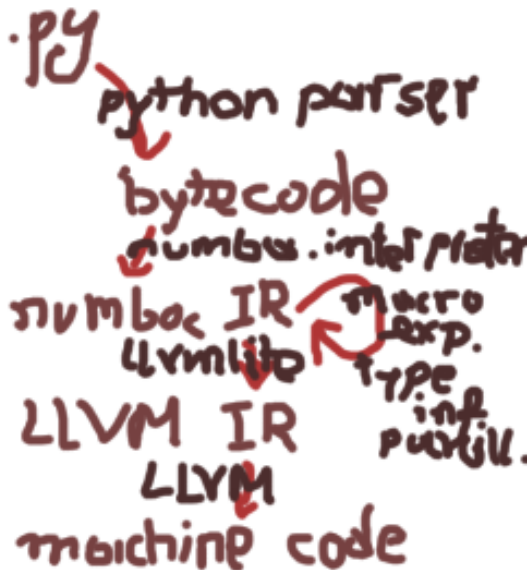
```
6
```

```
>>> import numpy as np
>>> x = np.eye(3)
>>> print('x:', x)
x: [[1. 0. 0.]
     [0. 1. 0.]
     [0. 0. 1.]]
>>> print('f(x):', f(x))
f(x): [[6. 0. 0.]
        [0. 6. 0.]
        [0. 0. 6.]]
```

```
>>> import numpy as np
>>> x = np.eye(3)
>>> print('x:', x)
x: [[1. 0. 0.]
     [0. 1. 0.]
     [0. 0. 1.]]
>>> print('f(x):', f(x))
f(x): [[6. 0. 0.]
        [0. 6. 0.]
        [0. 0. 6.]]
>>> f('abc')
'abcabcabcabcabc'
```

```
>>> f
CPUDispatcher(<function f at 0x...>)
>>> f.signatures
[(int64,),
 (array(float64, 2d, C),),
 (str,)]
>>> f.nopython_signatures
[(int64,) -> int64,
 (array(float64, 2d, C),) -> array(float64, 2d, C)]
```

Numba architecture



Some timings

```
def runningsum_loop(N):  
    s = 0  
    for i in range(N + 1):  
        s += i  
    return s
```

```
% timeit runningsum_loop(1_000_000)
```

54 ms \pm 973 μ s per loop (mean \pm std. dev. of 7 runs,
10 loops each)

```
def runningsum_list(N):  
    return sum([i for i in range(N + 1)])  
  
%timeit runningsum_list(1_000_000)  
52.7 ms ± 543 µs per loop (mean ± std. dev. of 7 runs,  
10 loops each)
```



```
def runningsum_list(N):  
    return sum([i for i in range(N + 1)])  
  
%timeit runningsum_list(1_000_000)  
52.7 ms  $\pm$  543  $\mu$ s per loop (mean  $\pm$  std. dev. of 7 runs,  
10 loops each)  
  
def runningsum_generator(N):  
    return sum(i for i in range(N + 1))  
  
%timeit runningsum_generator(1_000_000)  
50.4 ms  $\pm$  1.32 ms per loop (mean  $\pm$  std. dev. of 7 runs,  
10 loops each)
```

```
def runningsum_numpy(N):  
    return np.arange(N+1).sum()
```

```
%timeit runningsum_numpy(1_000_000)
```

1.41 ms \pm 76.8 μ s per loop (mean \pm std. dev. of 7 runs,
1000 loops each)

```
import numba

@numba.jit
def runningsum_numba_loop(N):
    s = 0
    for i in range(N + 1):
        s += i
    return s
```

```
import numba
```

```
@numba.jit
```

```
def runningsum_numba_loop(N):
```

```
    s = 0
```

```
    for i in range(N + 1):
```

```
        s += i
```

```
    return s
```

```
%timeit -r 1 -n 1 runningsum_numba_loop(1_000_000)
```

```
173 ms ± 0 ns per loop (mean ± std. dev. of 1 run,  
1 loop each)
```

```
import numba
```

```
@numba.jit
```

```
def runningsum_numba_loop(N):
```

```
    s = 0
```

```
    for i in range(N + 1):
```

```
        s += i
```

```
    return s
```

```
%timeit -r 1 -n 1 runningsum_numba_loop(1_000_000)
```

```
173 ms ± 0 ns per loop (mean ± std. dev. of 1 run,  
1 loop each)
```

```
%timeit runningsum_numba_loop(1_000_000)
```

```
195 ns ± 5.48 ns per loop (mean ± std. dev. of 7 runs,  
10_000_000 loops each)
```

```
long int sum(int N) {  
    long int s = 0;  
    for (int i=0; i <= N; i++)  
        s += i;  
    return s;  
}  
  
int main(int argc, char **argv) {  
    long int s;  
    for (int i = 0; i < 10000; i++)  
        s = sum(1000000);  
    printf("%ld\n", s);  
    return 0;  
}
```

gcc -g -Wall → 2.8 ms

gcc -g -Wall -O3 → <1 μs (naive)

gcc -g -Wall -O3 → 237 μs (external compilation unit)

Some timings — summary

	time / ms	
runningsum_loop	54	
runningsum_list	53	
runningsum_generator	50	
runningsum_numpy	1.41	
runningsum_numba_loop	173	(single iteration)
runningsum_numba_loop	0.000195	(repeated)
runningsum_c	2.8	-00
runningsum_c	<0.001	-03
runningsum_c	0.237	-03,
		seperate compilation units

Other numba features

Automatic parallelization

```
def trig_ident_np(x):  
    return (np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2).sum()/4
```

Automatic parallelization

```
def trig_ident_np(x):  
    return (np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2).sum()/4  
  
trig_ident_jit = numba.jit(trig_ident_np)  
trig_ident_jitp = numba.jit(parallel=True)(trig_ident_np)
```

Automatic parallelization

```
def trig_ident_np(x):  
    return (np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2 +  
            np.sin(x)**2 + np.cos(x)**2).sum()/4  
  
trig_ident_jit = numba.jit(trig_ident_np)  
trig_ident_jitp = numba.jit(parallel=True)(trig_ident_np)  
  
x = np.random.randn(500, 50_000)  
  
%timeit trig_ident_np(x)  
4.52 s ± 160 ms per loop  
  
%timeit trig_ident_jit(x)  
788 ms ± 24 ms per loop  
  
%timeit trig_ident_jitp(x)  
290 ms ± 7.38 ms per loop
```

Hardware support

- ▶ vector instructions (when CPU supports SSE, AVX, or AVX-512)
- ▶ Nvidia CUDA backend

Evaluation of numba

- ▶ good: native syntax and seamless integration
- ▶ good: excellent speed (when it works)
- ▶ bad: requires the whole LLVM backend to be present
- ▶ bad: hard to debug
- ▶ bad: not “reproducible”

Where is this all going?

Is Python a statically typed language?

Is Python a statically typed language?

```
# adder.py
```

```
def add(a:int, b:int = 1) -> int:  
    return a + b
```

```
add(1, 2)
```

```
add(1.2, 2.2)
```


Is Python a statically typed language?

```
# adder.py
```

```
def add(a:int, b:int = 1) -> int:  
    return a + b
```

```
add(1, 2)
```

```
add(1.2, 2.2)
```

```
$ mypy adder.py
```

```
adder.py:6: error: Argument 1 to "add" has incompatible  
type "float"; expected "int"
```

```
adder.py:6: error: Argument 2 to "add" has incompatible  
type "float"; expected "int"
```

The future?

- ▶ Python continues to be used a glue language
- ▶ Python code is seamlessly compiled with various backends
- ▶ Type hints are used where automatic type inference is insufficient

The End

Inspecting numba outputs

```
print(runningsum_numba_loop.inspect_llvm()[(numba.int64,)])
; ModuleID = 'runningsum_numba_loop'
source_filename = "<string>"
target datalayout = "e-m:e-i64:64-f80:128-n8:16:32:64-S128"
target triple = "x86_64-unknown-linux-gnu"

@_ZN08NumbaEnv8__main__25runningsum_numba_loop$242Ex" = common local_unnamed_addr @global i8* null
@.const.runningsum_numba_loop = internal constant [22 x i8] c"runningsum_numba_loop\00"
@PyExc_RuntimeError = external global i8
@".const.missing Environment" = internal constant [20 x i8] c"missing Environment\00"

; Function Attrs: norecurse nounwind
define i32 @"_ZN8__main__25runningsum_numba_loop$242Ex"(i64* noalias nocapture %retptr, { i8*, i32 }** noalias nocapture %args) {
entry:
    %.27 = add nsw i64 %arg.N, 1
    %.76 = icmp slt i64 %arg.N, 0
    %spec.select = select i1 %.76, i64 0, i64 %.27
    %.1214 = icmp sgt i64 %spec.select, 0
    br i1 %.1214, label %B20.lr.ph, label %B32

B20.lr.ph:
    %0 = xor i64 %spec.select, -1
    %1 = icmp sgt i64 %0, -2
    %smax = select i1 %1, i64 %0, i64 -2
    %2 = add i64 %spec.select, %smax
    %3 = add i64 %2, 1
    %4 = zext i64 %3 to i65
    %5 = zext i64 %2 to i65
    %6 = mul i65 %4, %5
    %7 = lshr i65 %6, 1
    %8 = trunc i65 %7 to i64
    %9 = add i64 %2, %8
    %10 = add i64 %9, 1
    br label %B32

; preds = %entry
```

Inspecting numba outputs

```
print(runningsum_numba_loop.inspect_asm()[(numba.int64,)])
```

```
.text
.file      "<string>"
.globl     _ZN8__main__25runningsum_numba_loop$242Ex
.p2align   4, 0x90
.type      _ZN8__main__25runningsum_numba_loop$242Ex,@function
_ZN8__main__25runningsum_numba_loop$242Ex:
    xorl    %ecx, %ecx
    testq   %rdx, %rdx
    leaq    1(%rdx), %rax
    cmovsq   %rcx, %rax
    testq   %rax, %rax
    jle     .LBB0_2
    movq     %rax, %rcx
    notq     %rcx
    cmpq     $-3, %rcx
    movq     $-2, %rdx
    cmovgq   %rcx, %rdx
    leaq     (%rax,%rdx), %rcx
    addq     %rax, %rdx
    addq     $1, %rdx
    mulxq    %rcx, %rax, %rdx
    shldq    $63, %rax, %rdx
    addq     %rdx, %rcx
    addq     $1, %rcx
.LBB0_2:
    movq     %rcx, (%rdi)
    xorl     %eax, %eax
    retq
.Lfunc_end0:
.size       _ZN8__main__25runningsum_numba_loop$242Ex, .Lfunc_end0-_ZN8__main__25runningsum_numba_loop$242Ex
.globl     _ZN7cpython8__main__25runningsum_numba_loop$242Ex
.p2align   4, 0x90
```